test accuracy (%) · cost relative to always-large

	GSM8K (n=20)	ARC-C (n=20)	MBPP (n=40)	MATH L4–5 (n=30)	MMLU-Pro (n=40)	Agent bug-fix (n=19)	Agent triage (n=20)
always-small	95 1% cost	85 1% cost	82 1% cost	53 2% cost	72 1% cost	68 0% cost	60 1% cost
always-medium	95 2% cost	100 1% cost	98 2% cost	70 3% cost	65 2% cost	79 2% cost	80 1% cost
always-large	90 100% cost	95 100% cost	90 100% cost	70 100% cost	85 100% cost	79 100% cost	80 100% cost
oracle	95 16% cost	100 1% cost	100 1% cost	83 41% cost	95 33% cost	89 31% cost	90 29% cost
zero-shot Laya	95 2% cost	95 1% cost	80 32% cost	67 21% cost	65 32% cost	74 2% cost	80 4% cost
SFT Laya	95 1% cost	85 1% cost	82 1% cost	53 2% cost	72 1% cost	68 0% cost	60 1% cost
RL auto-quality (λ=0)	90 100% cost	95 100% cost	88 79% cost	70 93% cost	85 85% cost	84 83% cost	70 77% cost
RL auto (λ=0.5)	95 2% cost	100 1% cost	95 2% cost	70 3% cost	65 2% cost	79 2% cost	75 1% cost